Learning to See in The Dark

What is the problem?

The primary challenge addressed in this paper is the difficulty of capturing high-quality images in extremely low-light conditions. It is quite challenging for the conventional cameras due to low photon count and low SNR. The images with short exposure suffer from noise, while the ones with long exposure can induce blur and is often impractical. While several denoising and enhancement techniques exist, they fail under extreme conditions, such as video-rate imaging at night with very little illumination. To tackle these issues, we introduced a machine learning based approach. We developed a pipeline for processing low-light images, based on end-to-end training of a fully convolutional network using a dataset of raw short exposure low-light images, with corresponding long-exposure reference image. The network operates directly on raw sensor data and replaces much of the traditional image processing pipeline, which tends to perform poorly on such data.

What has been done earlier?

Earlier work on low-light imaging includes various denoising techniques like BM3D, wavelet transforms, and deep learning models such as SSDA and convolutional networks, but these methods often relied on synthetic noise and struggled with real, extreme low-light data. Burst imaging, which combines multiple frames to reduce noise, is effective under moderate conditions but fails in extreme darkness due to alignment issues. Enhancement techniques like histogram equalization and Retinex models improve brightness but don't address severe noise.

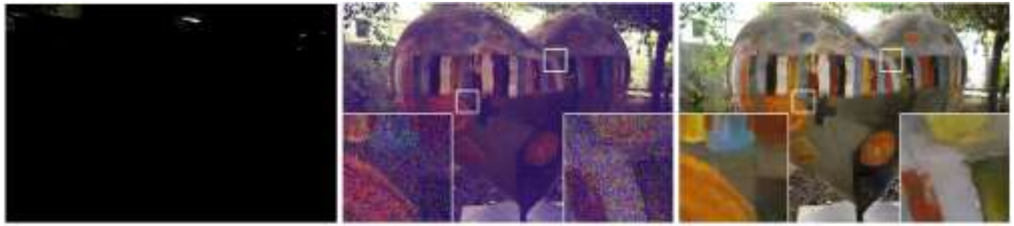

(a) JPEG image produced by camera (b) Raw data via traditional pipeline (c) Our result

Pulkit Sinha, B321059

What are the remaining challenges? What novel solution proposed by the authors to solve the problem?

The current solution does not address high dynamic range (HDR) tone mapping. The SID (See In the Dark) dataset is focused on static scenes, meaning the method has not yet been tested on dynamic scenes involving humans or moving objects, which would introduce further challenges such as motion blur and misalignment. The loss of details still occur in extreme low-light conditions, especially in the more challenging subsets of the dataset which can be improved further. Another limitation of the presented pipeline is that the amplification ratio must be chosen externally. It would be useful to infer a good amplification ratio from the input, akin to Auto ISO. Additionally, it has been assumed that a dedicated network is trained for a given camera sensor. The preliminary experiments with cross-sensor generalization are encouraging, and future work could further study the generalization abilities of low-light imaging networks. The presented pipeline takes 0.38 and 0.66 seconds to process full-resolution Sony and Fuji images, respectively; this is not fast enough for real-time processing at full resolution.

The authors proposed certain novel solutions to yield further improvements in image quality in the future work by redefining and systematically optimizing the network architecture and training procedure. The authors have also been looking to further improve the SID (See In the Dark) dataset to produce better and improved results in the future pipelines.

Pulkit Sinha, B321059